

Stable and Identifiable Bias Directions — Complete Textbook (12 Chapters)

Full Chinese textbook: TEXTBOOK_Stability_Identifiability.md | Chapters: reading_pack/md/*.md

Ch1 Linear Algebra: Directions, Projections, Cosine

1.1

1.1

Let $\theta \in \mathbb{D}$, $\|\theta\|=1$. Let x be a vector. The cosine of the angle between x and θ is

$$\text{score} = \theta^T x = \|x\| \cdot \cos(x, \theta)$$

Park's sparse linear combination $\sum c_i f_i$ where $f_i = \theta_i$

1

$$\cos(\theta_i, \theta_j) = \theta_i^T \theta_j / (\|\theta_i\| \cdot \|\theta_j\|) \in [-1, 1]$$

- Let $\theta^0, \theta^1, \dots, \theta^{20}$ be 20 bootstrap samples.
- Let $\text{Stab_boot} = E[\cos(\theta^i, \theta^j)]$

2

Let $w = w/\|w\|$. Let $\alpha = \theta^T w$. Let $P_\theta(x) = (\theta^T x) \theta$

3 Ch6

LoRA $W + \Delta W$, $\Delta W = b \cdot a^T$. Wang's $W + W \cdot b \cdot a^T$. Wang's W head

**Elhage p.1-6 Park p.4-6 (2)

1.2

- Let θ_i, θ_j be vectors. $E[\cos] = 0$, $\text{Var} \approx 1/D$. $D=896$. ~ 0.03 to 0.979
- Let $\mu_i = [1, 0]$, $\mu_j = [-1, 0]$. Let $\theta = \mu_i - \mu_j = [2, 0]$. Let $\theta = [1, 0]$. Let $|T|$ be mass-mean

Ch2 Transformer Anatomy: Residual Stream & Attention Head

2 Transformer Residual Stream Attention Head

2.1

```
r = Embed(tokens)
for l=1..L:
  o_{l,h} = Attn_{l,h}(r_{l-1}) ∈ ℝ^d // head
  o_l = concat_h o_{l,h} ∈ ℝ^{D·H}
  r_{l+1} = r_{l-1} + W_l · o_l + MLP(...) // Residual Stream
r_L → Unembed → logits
```

• `r`` Residual Stream head+MLP MVP `features[:,l,:]`

• `x_{l,h}=o_{l,h}` Attention Head Kim 1024

**`r` middle layer Kim Fig.2 `best layer 24` vs Kim `layer 16` vs **

** Elhage p.12-14 Fig.1 Kim Sec.2 (1)-(3) p.2-4

2.2

• Elhage Fig.1 residual head `W_{l,h} ∈ ℝ^{D×d}` `W_l = [W_{l,1}...W_{l,H}]`

Ch3 Probing: Correlation vs Causation Trap

3

3.1

- 1. $D=\{(x,y)\} \quad y=0 \text{ male, } 1 \text{ female}$
- 2. $\min_{\theta} \sum \text{CE}(y, \sigma(\theta x)) + \lambda ||\theta||^2 \quad \theta = \text{coef}$
- 3. $\text{Stab} < 0.3$

3.2 Hewitt & Liang

$y \neq \text{shuffled null}$

$\text{Stab} < 0.3$

3.3 vs

Nanda Othello 1.7% 20% Park non-linear probe $\leq$ linear probe LRH

Hewitt p.4-5 Control Task Nanda Othello-GPT

Ch4 Linear Representation Hypothesis: Why Linear?

4 LRH

4.1 Park

```
**`c` `f_c`  
x = Σ c_i · f_i + noise, c_i  
p(next token | x) ∝ exp( u_v · r_L )  
`f_c` `c` counterfactual pair `James/Mary`  
**  
1. **Log-Odds` `softmax CE` `logit` log-odds log-odds  
2. ** GD `f` Park Theorem 1-2
```

4.2

Park `linear probe ≈ non-linear probe` `steering` LRH superposition

** Park p.1-3, p.4-6 Theorem Gurnee p.2-4

Ch5 Kim's Method: Political Linear Probe & Steering

5 Kim Steer

5.1

- ** prompt → DW-NOMINATE $y \in [-1, +1]$
- $x_{\{l, h\}}$ Ridge $\min \Sigma (y - \theta x)^2 + \lambda ||\theta||^2$ Spearman Top-K=32 head

5.2

- Fig.2 Top Head Middle Layer (≈ 16) $p \approx 0.7$
- $x^{\alpha}_{\{l, h\}} = x_{\{l, h\}} + \alpha \cdot \sigma_{\{l, h\}} \cdot \theta_{\{l, h\}}$ head $\alpha \in [-3, 3]$
- Fig.4 α

5.3

Kim θ $\alpha \cdot \sigma \cdot \theta$ Wang

** Kim p.2-4 Sec.2, p.4-8 Sec.3, p.11-14 Sec.6 (8) Fig.4

Ch6 Wang's Critique: Optimal Editing & Null Distribution

6 Wang

6.1 ITI

ITI $\theta = \mu - \mu$ heuristic Wang

6.2

head TruthfulQA

TruthfulQA (x, y, y) **IPO**

$$\max \sum [\log(\pi_\phi(y|x)/\pi(y|x)) / \pi_\phi(y|x)/\pi(y|x)) - \tau^{1/2}]^2$$

LoRA

1. **LoRA** $r_{\text{LoRA}} = r + (W + b a) o = r_{\text{orig}} + a, o \cdot b$

• b W **head**

2. **Wang** $b \rightarrow W \cdot b$

...

$r_{\text{reparam}} = r_{\text{orig}} + a, o \cdot W \cdot b = r_{\text{orig}} + a, o \cdot \sum_h W_{\{h\}} \cdot b_{\{h\}}$

...

• $b_{\{h\}}$ $\theta_{\{h\}}$ a, o α

• $b_{\{h\}} = 0$ **head** **head**

1. Top-16 $\approx$ Full-model (Fig.2) —

2. **Random-16** $\approx$ Top-16 (Fig.3) —

3. **Head** $\approx$ Full-model **head** (Fig.4) —

4. **heuristic** **Top** $\gg$ **Random** **optimal** **Top** $\approx$ **Random**

Wang p.3-4 Sec.2, p.6-9 Sec.4-6 Fig.2-4 (7)(8)

6.3

• **Fig.1** **Heuristic** **Top** **Info*Truth** **Random** $p = 1.6e-8$ —

• **Fig.3** **Optimal** **Random** **Top** —

• **Fig.4a** **24** **head** **5** $\approx$ **Top-16**

Ch7 Bias Measurement: Behavior to Representation

7 ■■■■■■■■■■■■

7.1 ■■■■

- **WinoBias** `The nurse said she...` he/she
- **OccuGender** 1) 2) BLS 3) 4) 5) `At work, {name} the {occupation}...` leave-one-template-out 1

7.2 ■■■■

≠

- **+ logit Δ + LEACE θ**

OccuGender p.2-3 Desiderata, p.4 Fig.2 Zhao WinoBias p.2

Ch8 Diagnostic Design: Four-Control Dual-Track

8 ResearchStudio C03

8.1 ResearchStudio C03

- 1. **James vs Mary`
- 2. **`
- 3. **`20xrandom` + `shuffled` + `full-model`

8.2 C01

` Wang Fig.3 Kim Fig.2

8.3 PrivGap

D_true → θ_{true} → Heuristic ΔM_{true}
D_shuf → θ_{shuf} → Heuristic ΔM_{shuf} → PrivGap_heu
20xrand → θ_{rand} → Heuristic ΔM_{rand} → Optimal ΔM^* (LoRA-IPO) → PrivGap_opt

Ch9 Statistical Testing: Bootstrap & Permutation

9 Bootstrap Permutation

9.1 Bootstrap Stability

```
for b=1..20:
  sample 80% male + 80% female with replacement
   $\theta^b = \text{probe}(X_b, y_b)$ 
  Stab = mean_{i≠j} cos( $\theta^i, \theta^j$ )
```

- ■■■■■■■■male/female ■■■■■■
- ■■ `0.979` ■ `mean off-diag` ■■■■■■ `std` ■ `shuffled null` ■

9.2 Permutation ■■ PrivGap

```
null = { PrivGap_rand^1 ... PrivGap_rand^20 }
p = (# null ≥ PrivGap_true + 1)/(20+1)
```

- $p < 0.01$ ■■■■■■
- ■■■■ $\text{corr}(\text{Stab}, \text{PrivGap})$ ■■■■ Stability ■■■■ Identifiability ■

****■■■■■** Wasserman Ch.8 Bootstrap ■ Efron 1979**

...

100

- ```

1. ██
2. ████████████████████████
3. ███ Ch6 ████████`████ Wang Fig.3 █████ Kim Fig.2` █████`Fig.2 heuristic████Fig.3 optimal████`██
4. ███████`DEFINITIONS.md` █ `IDEA_CARD.md` ████████████████

███████ `reading_pack/` █ `A1_A2...` █████ PDF███████

```

## Overview: What Are We Answering?

■■■■■■■■■■ — ■■

# Stable and Identifiable Bias Directions in LLMs: From Linear Representation to Privileged Causality

## ■ Toward Stable and Identifiable Bias Directions

[illegible]

...

0

**\*\*■■■■■Research Question■■■\*\***

[illegible]

**\*\*■■■■■ VS ■■■■■\*\***

- \*\*■■■1 Kim 2503.02080■■■`x+ασθ` steer■■■
- \*\*■■■2 Park LRH■■■
- \*\*■■■ Wang 2502.11447■■■

\*\*☐☐☐☐☐\*\*

Ch1 ■■■■ → Ch2 Transformer■■ → Ch3 ■■ → Ch4 LRH■■ → Ch5 Kim■■■ → Ch6 Wang■■■ → Ch7 ■■■■ → Ch8 ■■■■ → Ch9 ■■■■

— — —

# Appendix A: Page Map & Glossary

## ■ ■ ■ A ■ ■ ■ ■ ■ ■ ■ ■ ■ ■

|     |                             |                    |                     |  |
|-----|-----------------------------|--------------------|---------------------|--|
| ■ ■ | ■ ■ ■ ■ ■                   | ■ ■ ■ ■ ■          | ■ ■ ■ ■ ■           |  |
| Ch1 | Elhage Transformer Circuits | p.1-6, p.12-14     | Fig.1 Residual      |  |
| Ch1 | Park 2403.03867             | p.4-6 ■ ■ (2)      | -                   |  |
| Ch2 | Elhage                      | p.12-14            | Fig.1               |  |
| Ch2 | Kim 2503.02080 Sec.2        | p.2-4 ■ ■ (1)-(3)  | -                   |  |
| Ch3 | Hewitt & Liang 2019         | p.4-5 Control Task | Fig.2               |  |
| Ch4 | Park                        | p.1-3, p.4-6       | Fig.2, Fig.4        |  |
| Ch5 | Kim 2503.02080              | p.1-14             | Fig.1, Fig.2, Fig.4 |  |
| Ch6 | Wang 2502.11447             | p.1-10, ■ ■ (7)(8) | Fig.1,3,4           |  |
| Ch7 | OccuGender / WinoBias       | OccuGender p.2-4   | Fig.2               |  |
| Ch8 | ResearchStudio C03/C01      | -                  | ■ ■ ■ ■             |  |
| Ch9 | Wasserman Ch.8              | Ch.8               | -                   |  |

## ■ ■ ■ B ■ ■ ■ ■ ■

Stability / Identifiability / PrivGap / ITI / LoRA / IPO / Residual Stream / Attention Head / Bootstrap / Permutation

# Appendix B: 10-Day Checklist

## ■ ■ C 10 ■ ■ ■ ■ ■ ■ ■ ■ ■ ■

|     |                            |                           |                         |
|-----|----------------------------|---------------------------|-------------------------|
| Day | ■ ■                        | ■ ■                       | ■ ■                     |
| 1-2 | Ch1+Ch2                    | ■ ■ ■ ■ Residual Stream + | ■ ■ ■ ■   ■             |
| 3-4 | Ch3+Ch4                    | ■ ■ Park Theorem1         | ■                       |
| 5-7 | Ch5 vs Ch6                 | ■ ■ ■                     | ■ Kim vs Wang ■ ■ ■   ■ |
| 8   | Ch7                        | ■ ■ shuffled null <0.3    | ■                       |
| 9   | Ch8+Ch9                    | ■ Related Work            | ■ ■ ■   ■               |
| 10  | ■ run_bias_direction_v2.py | ■ Fig1                    | ■ ■ ■ ■   ■             |

\*\* ■ ■ ■ ■ ■ ■ \*\* ■ ■ ■ ■ ` ■ ■ ■ ■ Wang Fig.3 ■ ■ ■ ■ Kim Fig.2` ■ ■ ■ `Fig.2 heuristic ■ ■ ■ ■ Fig.3 optimal` ■ ■ ■ ■